



# Face Emotion Classification


Ravi shankar Purushothaman

# What we have done so far:

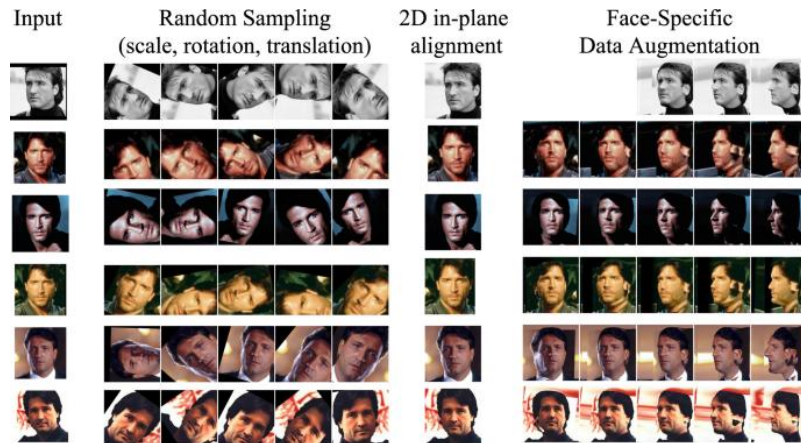


- Building a CNN: A Convolutional Neural Network (CNN) from scratch to classify facial expressions into predefined emotion categories.
- RESNET50 Model: Implementing the RESNET50 architecture, pre-trained on ImageNet, to apply transfer learning for emotion recognition tasks.
- EfficientNetV2L Models: Implementing the EfficientNetV2L architecture, a state-of-the-art convolutional neural network known for its efficiency and effectiveness in image classification tasks. Fine-tuning model to improve the overall accuracy of val.

# Data Preprocessing:



## Data Generators for Training: Data Augmentation



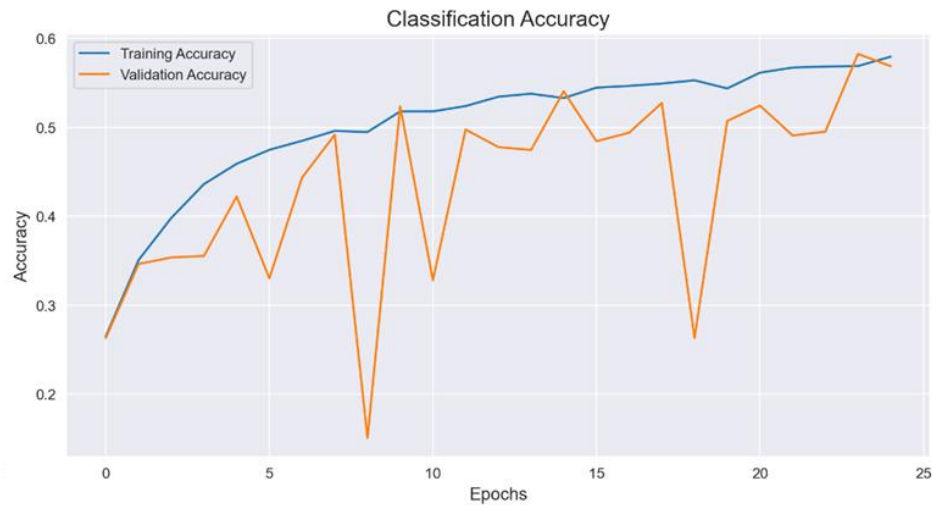
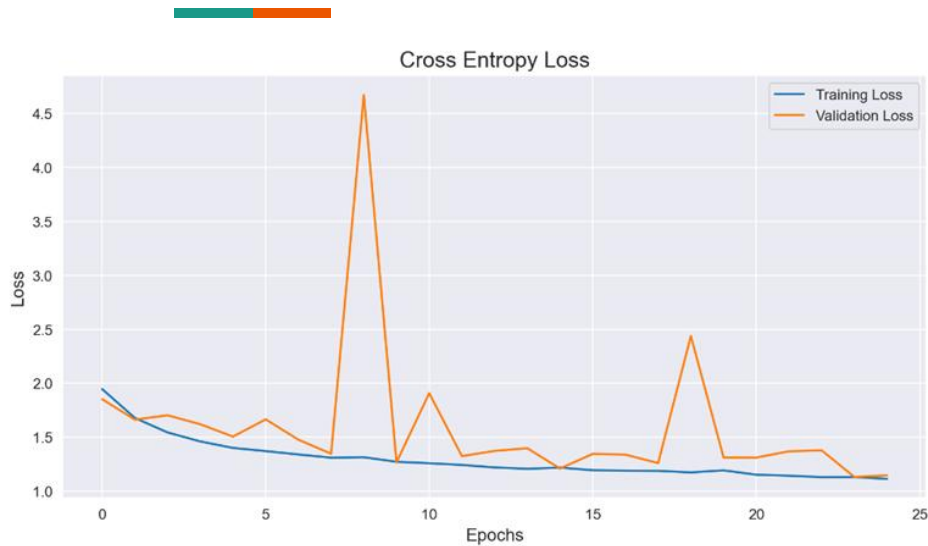
- The dataset used in this project consists of facial images annotated with emotion labels.
- Each image is categorized into one of several emotion classes, including 'happy', 'sad', 'angry', 'surprise', 'fear', 'disgust', and 'neutral'.
- The dataset is preprocessed to standardize image sizes and formats, ensuring compatibility with the chosen deep-learning models.

- CNN from Scratch



```
def define_model(image_size, image_channel):  
    """  
    Define a Convolutional Neural Network (CNN) model for image classification.  
    """  
    model = Sequential()  
  
    # Input Layer  
    model.add(Input(shape=(image_size[0], image_size[1], image_channel))) # Input layer to specify  
  
    # Convolutional Layers  
    model.add(Conv2D(32, (3, 3), activation='relu'))  
    model.add(BatchNormalization())  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Dropout(0.2))  
  
    model.add(Conv2D(64, (3, 3), activation='relu'))  
    model.add(BatchNormalization())  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Dropout(0.2))  
  
    model.add(Conv2D(128, (3, 3), activation='relu'))  
    model.add(BatchNormalization())  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Dropout(0.2))  
  
    model.add(Conv2D(256, (3, 3), activation='relu'))  
    model.add(BatchNormalization())  
    model.add(MaxPooling2D(pool_size=(2, 2)))  
    model.add(Dropout(0.2))  
  
    # Fully Connected Dense Layers  
    model.add(Flatten())  
    model.add(Dense(512, activation='relu'))  
    model.add(BatchNormalization())  
    model.add(Dropout(0.2))  
  
    # Output layer  
    model.add(Dense(NB_CLASSES, activation='softmax'))  
  
    model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

# CNN Training Results:



loss: 1.1122 - accuracy: 0.5796 - val\_loss: 1.1455 - val\_accuracy: 0.5688

# CNN Eval Results (Test):



True: happy  
Predicted: happy



True: happy  
Predicted: happy



True: neutral  
Predicted: neutral



True: neutral  
Predicted: neutral



True: neutral  
Predicted: neutral



True: neutral  
Predicted: sad



True: neutral  
Predicted: neutral



True: neutral  
Predicted: sad



True: neutral  
Predicted: neutral




True: neutral


True: neutral

True: neutral

# Classes are unbalanced?



# Classes are unbalanced?







# RESNET (50) - Applying Transfer Learning

Model: 'sequential'


Layer (type)	Output Shape	Param #
=====		
resnet50 (Functional)	(None, 2, 2, 2048)	23587712
global_max_pooling2d (GlobalMaxPooling2D)	(None, 2048)	0
dropout (Dropout)	(None, 2048)	0
dense (Dense)	(None, 7)	14343
=====		
Total params: 23,602,055		
Trainable params: 23,548,935		
Non-trainable params: 53,120		

# RESNET 50 Training Results:



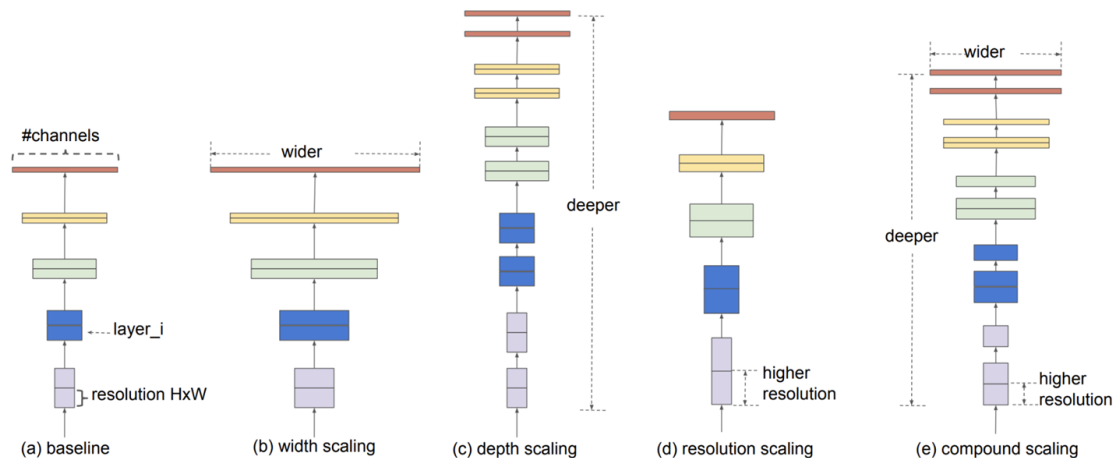
Loss: 0.1493, Accuracy: 0.8780 Validation Loss: 0.3230, Validation Accuracy: 0.6013

The current model achieved a categorical accuracy of 60.43%!



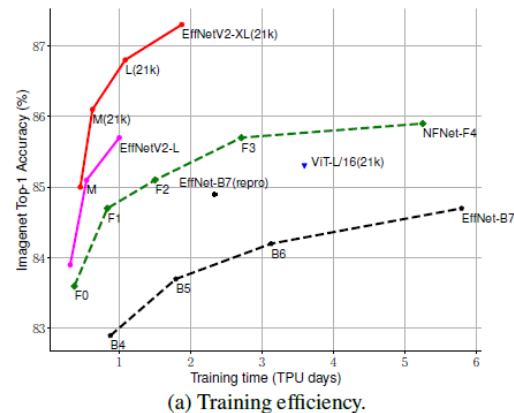

RESNET 50 Eval Results:

# EfficientNetV2L



**Figure 2. Model Scaling.** (a) is a baseline network example; (b)-(d) are conventional scaling that only increases one dimension of network width, depth, or resolution. (e) is our proposed compound scaling method that uniformly scales all three dimensions with a fixed ratio.

EfficientNetV2 is a type convolutional neural network that has faster training speed and better parameter efficiency than previous models. To develop these models, the authors use a combination of training-aware neural architecture search and scaling, to jointly optimize training speed.





# EfficientNetV2L - 1)

Model: "model\_3"



Layer (type)	Output Shape	Param #
input_layer (InputLayer)	[(None, 224, 224, 3)]	0
efficientnetv2-l (Functional)	(None, None, None, 1280)	117746848
global_average_pooling (GlobalAveragePooling2D)	(None, 1280)	0
Output_layer (Dense)	(None, 7)	8967

=====  
Total params: 117755815 (449.20 MB)  
Trainable params: 8967 (35.03 KB)  
Non-trainable params: 117746848 (449.17 MB)

loss: 1.1869 - accuracy: 0.5553 - val\_loss: 1.2284  
- val\_accuracy: 0.5400




# EfficientNetV2L - Unfreeze Last 20 Layers



accuracy: 0.8724 - val\_loss: 1.2964 - val\_accuracy: 0.6488

## EfficientNetV2L - Unfreeze Last 45 Layers



loss: 0.3212 - accuracy: 0.8990 - val\_loss: 1.0750 - val\_accuracy: 0.6622

# Model Performance Comparison & Conclusion

## CNN

Loss: 1.1122, Accuracy: 0.5796

Validation Loss: 1.1455, Validation Accuracy: 0.5688

## RESNET50

Loss: 0.1493, Accuracy: 0.8780

Validation Loss: 0.3230, Validation Accuracy: 0.6013

## Base EfficientNetV2L

Loss: 1.1869, Accuracy: 0.5553

Validation Loss: 1.2284, Validation Accuracy: 0.5400

## EfficientNetV2L (Unfreeze last 20 layers)

Loss: 0.3642, Accuracy: 0.8724

Validation Loss: 1.2964, Validation Accuracy: 0.6488

## Efficient V2L (Unfreeze last 45 layers)

Loss: 0.3212, Accuracy: 0.8990

Validation Loss: 1.0750, Validation Accuracy: 0.6622

- Transitioning from a basic CNN to RESNET50 significantly enhanced accuracy and validation metrics. Deeper architectures excel in feature extraction and representation learning.
- EfficientNetV2L, especially with deeper layer unfreezing, demonstrated robust performance gains. Its efficiency and scalability effectively capture intricate features crucial for emotion recognition.
- Regularization and Stability
- Techniques like dropout, batch normalization, and fine-tuning strategies mitigate overfitting and ensure model stability across architectures.